

## Lecture 8

# Concurrency control in applications

# Collaboration and conflict resolution

Nowadays we use a lot of **collaboration software**:

- ▶ **Examples:** calendar sync (last lecture), Google Docs, ...

# Collaboration and conflict resolution

Nowadays we use a lot of **collaboration software**:

- ▶ **Examples:** calendar sync (last lecture), Google Docs, ...
- ▶ Several users/devices working on a shared file/document
- ▶ Each user device has local replica of the data

# Collaboration and conflict resolution

Nowadays we use a lot of **collaboration software**:

- ▶ **Examples:** calendar sync (last lecture), Google Docs, ...
- ▶ Several users/devices working on a shared file/document
- ▶ Each user device has local replica of the data
- ▶ Update local replica anytime (even while offline), sync with others when network available

# Collaboration and conflict resolution

Nowadays we use a lot of **collaboration software**:

- ▶ **Examples:** calendar sync (last lecture), Google Docs, ...
- ▶ Several users/devices working on a shared file/document
- ▶ Each user device has local replica of the data
- ▶ Update local replica anytime (even while offline), sync with others when network available
- ▶ **Challenge:** how to reconcile concurrent updates?

# Collaboration and conflict resolution

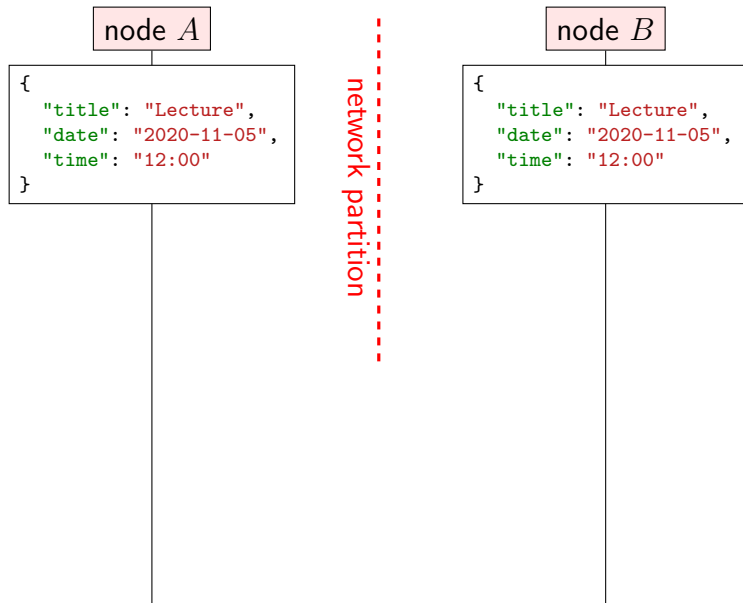
Nowadays we use a lot of **collaboration software**:

- ▶ **Examples:** calendar sync (last lecture), Google Docs, ...
- ▶ Several users/devices working on a shared file/document
- ▶ Each user device has local replica of the data
- ▶ Update local replica anytime (even while offline), sync with others when network available
- ▶ **Challenge:** how to reconcile concurrent updates?

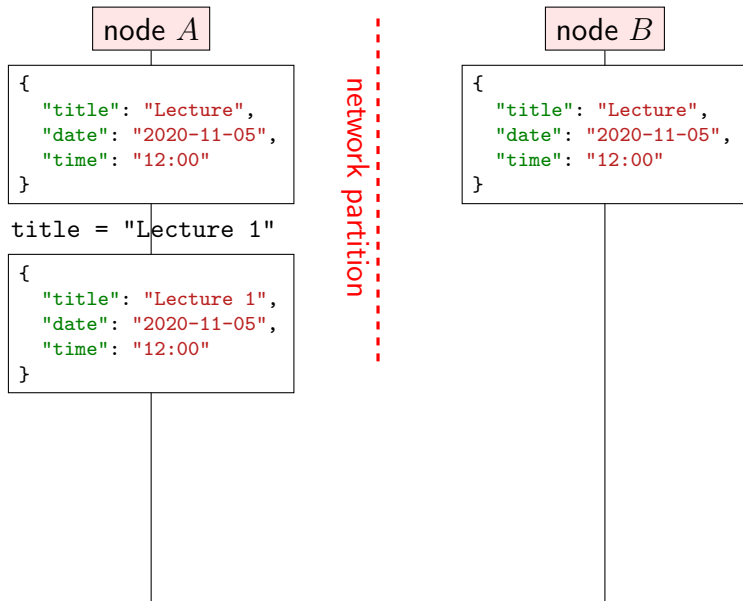
Families of **algorithms**:

- ▶ Conflict-free Replicated Data Types (**CRDTs**)
  - ▶ Operation-based
  - ▶ State-based
- ▶ Operational Transformation (**OT**)

# Conflicts due to concurrent updates

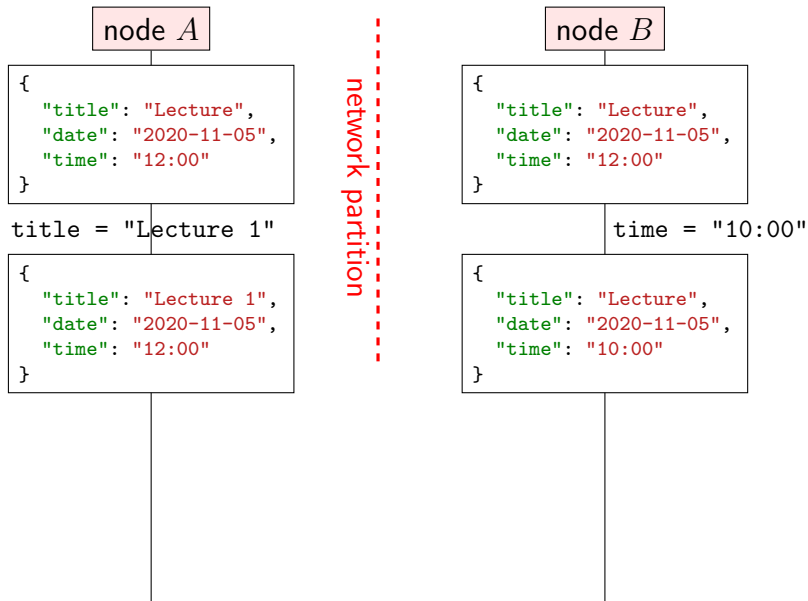


# Conflicts due to concurrent updates

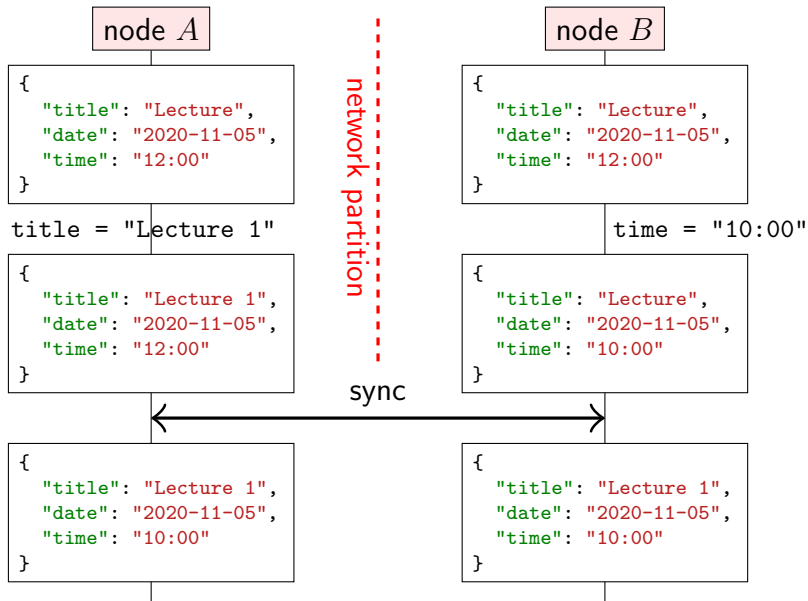




# Conflicts due to concurrent updates



# Conflicts due to concurrent updates



# Operation-based map CRDT

**on** initialisation **do**

$values := \{\}$

**end on**

**on** request to read value for key  $k$  **do**

**if**  $\exists t, v. (t, k, v) \in values$  **then return**  $v$  **else return** null

**end on**

**on** request to set key  $k$  to value  $v$  **do**

$t := \text{newTimestamp}()$   $\triangleright$  globally unique, e.g. Lamport timestamp

**broadcast** (set,  $t, k, v$ ) by reliable broadcast (including to self)

**end on**

**on** delivering (set,  $t, k, v$ ) by reliable broadcast **do**

$previous := \{(t', k', v') \in values \mid k' = k\}$

**if**  $previous = \{\} \vee \forall (t', k', v') \in previous. t' < t$  **then**

$values := (values \setminus previous) \cup \{(t, k, v)\}$

**end if**

**end on**

# Operation-based CRDTs

Reliable broadcast may deliver updates in any order:

- ▶ broadcast (set,  $t_1$ , "title", "Lecture 1")
- ▶ broadcast (set,  $t_2$ , "time", "10:00")

# Operation-based CRDTs

Reliable broadcast may deliver updates in any order:

- ▶ broadcast (set,  $t_1$ , "title", "Lecture 1")
- ▶ broadcast (set,  $t_2$ , "time", "10:00")

Recall **strong eventual consistency**:

- ▶ **Eventual delivery:** every update made to one non-faulty replica is eventually processed by every non-faulty replica.
- ▶ **Convergence:** any two replicas that have processed the same set of updates are in the same state

# Operation-based CRDTs

Reliable broadcast may deliver updates in any order:

- ▶ broadcast (set,  $t_1$ , "title", "Lecture 1")
- ▶ broadcast (set,  $t_2$ , "time", "10:00")

Recall **strong eventual consistency**:

- ▶ **Eventual delivery:** every update made to one non-faulty replica is eventually processed by every non-faulty replica.
- ▶ **Convergence:** any two replicas that have processed the same set of updates are in the same state

CRDT algorithm implements this:

- ▶ Reliable broadcast ensures every operation is eventually delivered to every (non-crashed) replica

# Operation-based CRDTs

Reliable broadcast may deliver updates in any order:

- ▶ broadcast (set,  $t_1$ , "title", "Lecture 1")
- ▶ broadcast (set,  $t_2$ , "time", "10:00")

Recall **strong eventual consistency**:

- ▶ **Eventual delivery**: every update made to one non-faulty replica is eventually processed by every non-faulty replica.
- ▶ **Convergence**: any two replicas that have processed the same set of updates are in the same state

CRDT algorithm implements this:

- ▶ Reliable broadcast ensures every operation is eventually delivered to every (non-crashed) replica
- ▶ Applying an operation is **commutative**: order of delivery doesn't matter

# State-based map CRDT

The operator  $\sqcup$  merges two states  $s_1$  and  $s_2$  as follows:

$$s_1 \sqcup s_2 = \{(t, k, v) \in (s_1 \cup s_2) \mid \nexists (t', k', v') \in (s_1 \cup s_2). k' = k \wedge t' > t\}$$

**on** initialisation **do**

$values := \{\}$

**end on**

**on** request to read value for key  $k$  **do**

**if**  $\exists t, v. (t, k, v) \in values$  **then return**  $v$  **else return** null

**end on**

**on** request to set key  $k$  to value  $v$  **do**

$t := \text{newTimestamp}()$   $\triangleright$  globally unique, e.g. Lamport timestamp

$values := \{(t', k', v') \in values \mid k' \neq k\} \cup \{(t, k, v)\}$

**broadcast**  $values$  by best-effort broadcast

**end on**

**on** delivering  $V$  by best-effort broadcast **do**

$values := values \sqcup V$

**end on**



# State-based CRDTs

Merge operator  $\sqcup$  must satisfy:  $\forall s_1, s_2, s_3 \dots$

- ▶ **Commutative:**  $s_1 \sqcup s_2 = s_2 \sqcup s_1$ .
- ▶ **Associative:**  $(s_1 \sqcup s_2) \sqcup s_3 = s_1 \sqcup (s_2 \sqcup s_3)$ .
- ▶ **Idempotent:**  $s_1 \sqcup s_1 = s_1$ .

# State-based CRDTs

Merge operator  $\sqcup$  must satisfy:  $\forall s_1, s_2, s_3 \dots$

- ▶ **Commutative:**  $s_1 \sqcup s_2 = s_2 \sqcup s_1$ .
- ▶ **Associative:**  $(s_1 \sqcup s_2) \sqcup s_3 = s_1 \sqcup (s_2 \sqcup s_3)$ .
- ▶ **Idempotent:**  $s_1 \sqcup s_1 = s_1$ .

State-based versus operation-based:

- ▶ Op-based CRDT typically has smaller messages
- ▶ State-based CRDT can tolerate message loss/duplication

# State-based CRDTs

Merge operator  $\sqcup$  must satisfy:  $\forall s_1, s_2, s_3 \dots$



- ▶ **Commutative:**  $s_1 \sqcup s_2 = s_2 \sqcup s_1$ .
- ▶ **Associative:**  $(s_1 \sqcup s_2) \sqcup s_3 = s_1 \sqcup (s_2 \sqcup s_3)$ .
- ▶ **Idempotent:**  $s_1 \sqcup s_1 = s_1$ .

State-based versus operation-based:

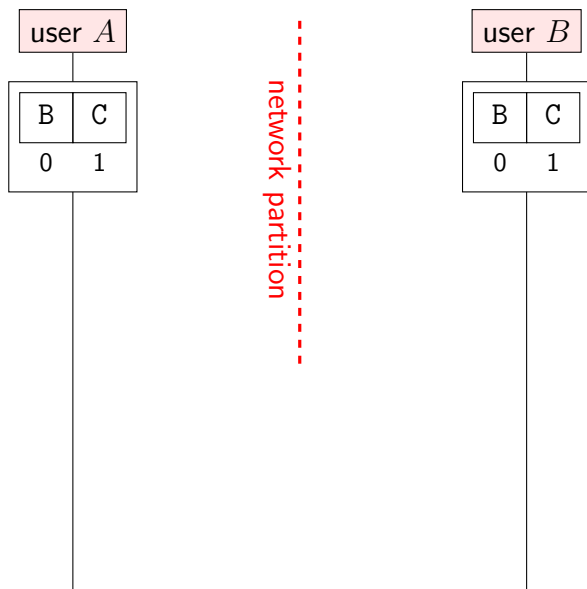
- ▶ Op-based CRDT typically has smaller messages
- ▶ State-based CRDT can tolerate message loss/duplication

Not necessarily uses broadcast:

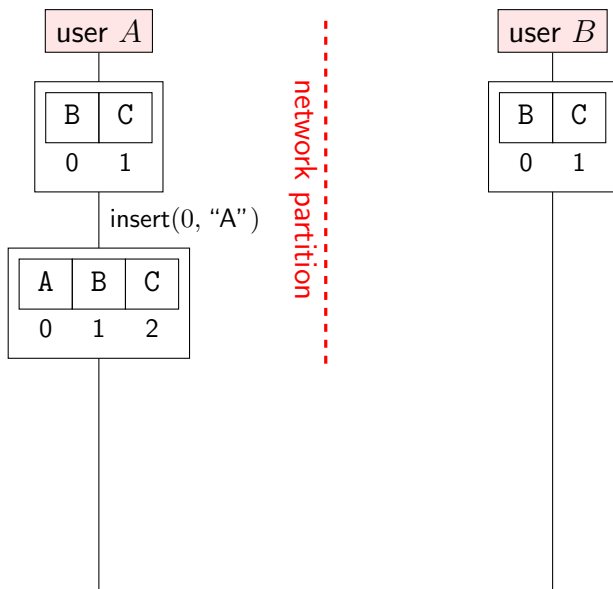
- ▶ Can also merge concurrent updates to replicas e.g. in quorum replication, anti-entropy, ...



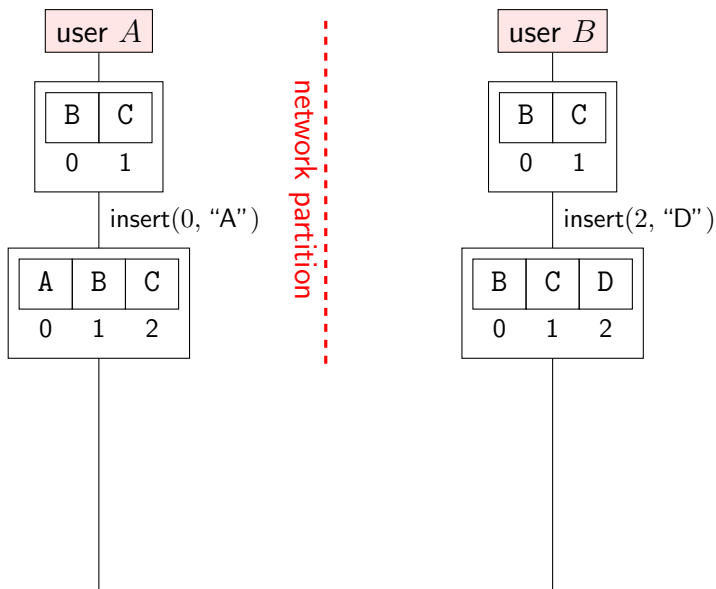
# Collaborative text editing: the problem



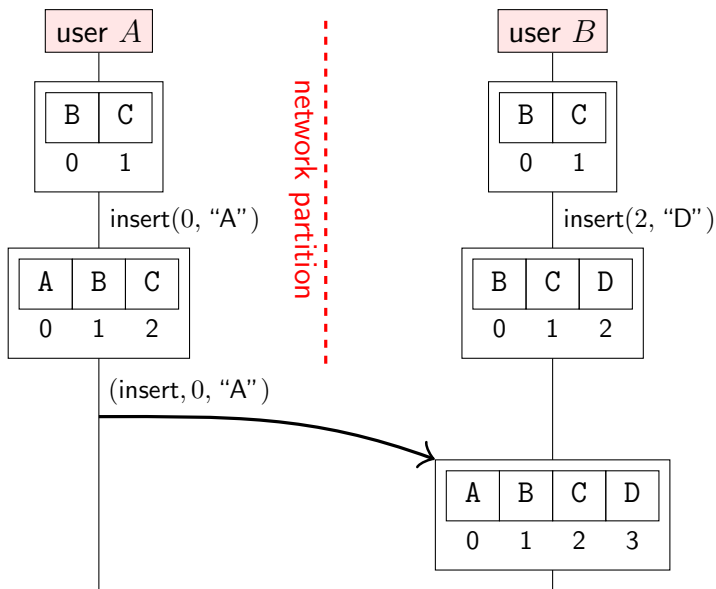
# Collaborative text editing: the problem



# Collaborative text editing: the problem

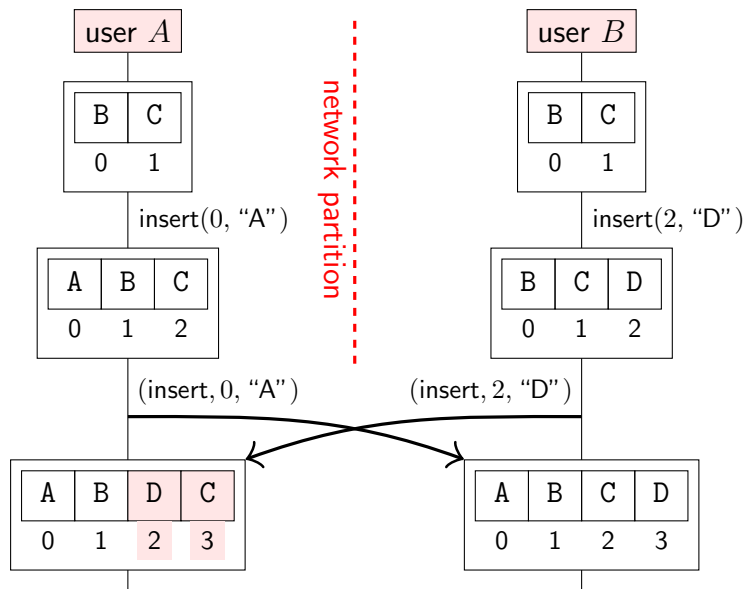


# Collaborative text editing: the problem







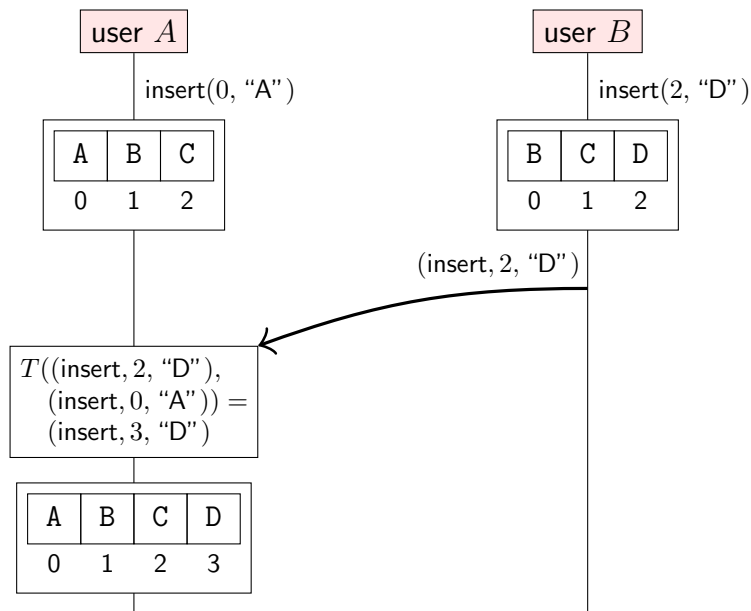
# Collaborative text editing: the problem



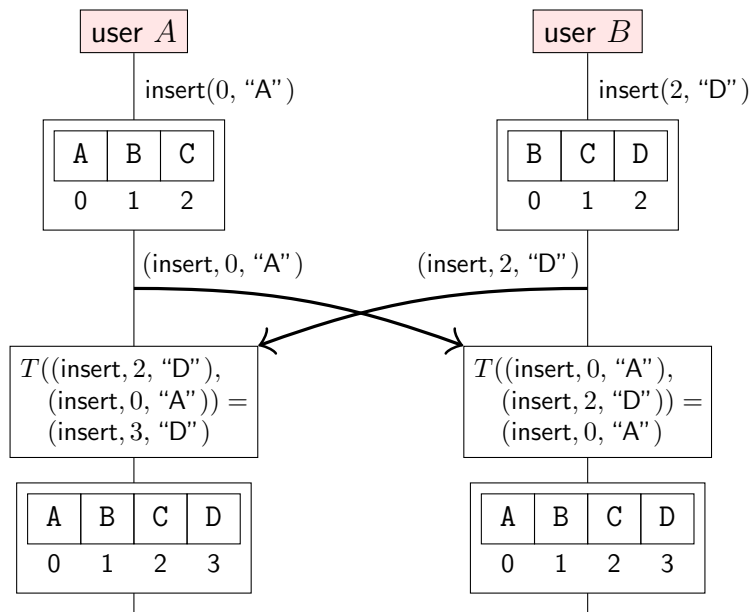
# Operational transformation





# Operational transformation




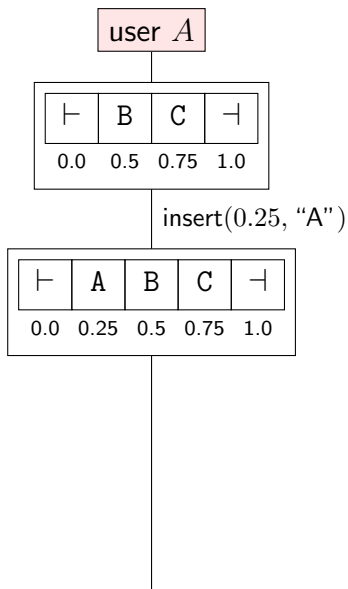
# Operational transformation




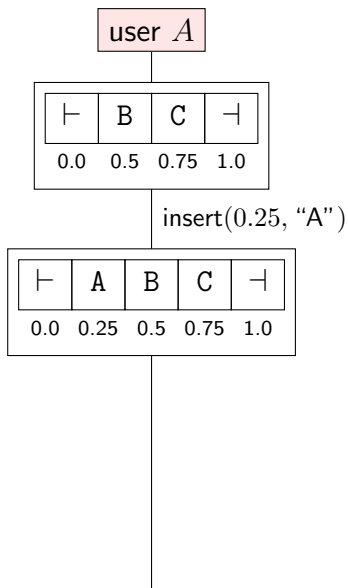
# Text editing CRDT



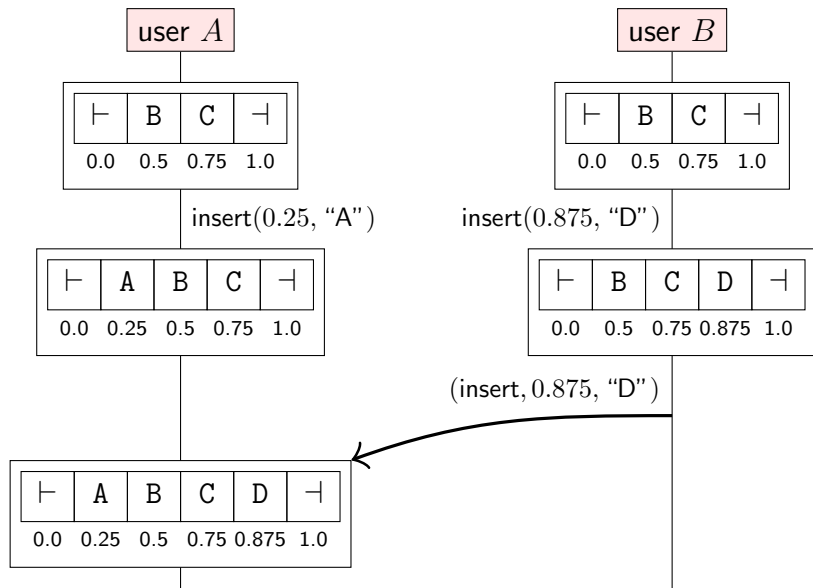
# Text editing CRDT



# Text editing CRDT

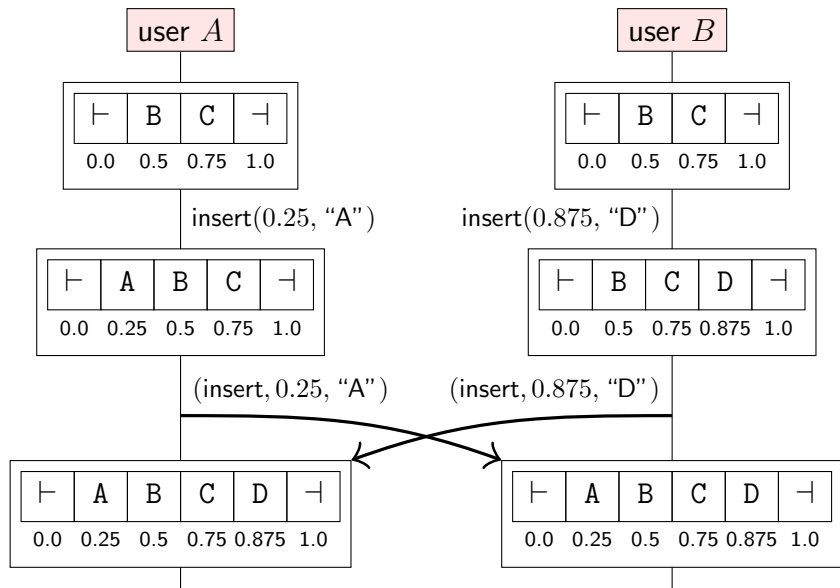


# Text editing CRDT





# Text editing CRDT



# Operation-based text CRDT (1/2)

**function** ELEMENTAT(*chars*, *index*)

*min* = the unique triple  $(p, n, v) \in \text{chars}$  such that

$\nexists (p', n', v') \in \text{chars}. p' < p \vee (p' = p \wedge n' < n)$

**if** *index* = 0 **then return** *min*

**else return** ELEMENTAT(*chars* \ {*min*}, *index* - 1)

**end function**

**on** initialisation **do**

*chars* := {(0, null, ⊢), (1, null, ⊣)}

**end on**

**on** request to read character at index *index* **do**

**let**  $(p, n, v) := \text{ELEMENTAT}(\text{chars}, \text{index} + 1)$ ; **return** *v*

**end on**

**on** request to insert character *v* at index *index* at node *nodeId* **do**

**let**  $(p_1, n_1, v_1) := \text{ELEMENTAT}(\text{chars}, \text{index})$

**let**  $(p_2, n_2, v_2) := \text{ELEMENTAT}(\text{chars}, \text{index} + 1)$

**broadcast** (insert,  $(p_1 + p_2)/2$ , *nodeId*, *v*) by causal broadcast

**end on**

# Operation-based text CRDT (2/2)

**on** delivering (insert,  $p, n, v$ ) by causal broadcast **do**  
     $chars := chars \cup \{(p, n, v)\}$   
**end on**

**on** request to delete character at index  $index$  **do**  
    **let**  $(p, n, v) := \text{ELEMENTAT}(chars, index + 1)$   
    **broadcast** (delete,  $p, n$ ) by causal broadcast  
**end on**

**on** delivering (delete,  $p, n$ ) by causal broadcast **do**  
     $chars := \{(p', n', v') \in chars \mid \neg(p' = p \wedge n' = n)\}$   
**end on**

- ▶ Use causal broadcast so that insertion of a character is delivered before its deletion
- ▶ Insertion and deletion of different characters commute

# Google's Spanner

A database system with millions of nodes, petabytes of data, distributed across datacenters worldwide

# Google's Spanner

A database system with millions of nodes, petabytes of data, distributed across datacenters worldwide

Consistency properties:

- ▶ **Serializable** transaction isolation
- ▶ **Linearizable** reads and writes

# Google's Spanner

A database system with millions of nodes, petabytes of data, distributed across datacenters worldwide

Consistency properties:

- ▶ **Serializable** transaction isolation
- ▶ **Linearizable** reads and writes
- ▶ Many **shards**, each holding a subset of the data; atomic commit of transactions across shards

# Google's Spanner

A database system with millions of nodes, petabytes of data, distributed across datacenters worldwide

Consistency properties:

- ▶ **Serializable** transaction isolation
- ▶ **Linearizable** reads and writes
- ▶ Many **shards**, each holding a subset of the data; atomic commit of transactions across shards

Many standard techniques:

- ▶ State machine replication (Paxos) within a shard
- ▶ Two-phase locking for serializability
- ▶ Two-phase commit for cross-shard atomicity

# Google's Spanner

A database system with millions of nodes, petabytes of data, distributed across datacenters worldwide

Consistency properties:

- ▶ **Serializable** transaction isolation
- ▶ **Linearizable** reads and writes
- ▶ Many **shards**, each holding a subset of the data; atomic commit of transactions across shards

Many standard techniques:

- ▶ State machine replication (Paxos) within a shard
- ▶ Two-phase locking for serializability
- ▶ Two-phase commit for cross-shard atomicity

The interesting bit: read-only transactions require **no locks!**



# Consistent snapshots

A read-only transaction observes a **consistent snapshot**:

If  $T_1 \rightarrow T_2$  (e.g.  $T_2$  reads data written by  $T_1$ )...

- ▶ Snapshot reflecting writes by  $T_2$  also reflects writes by  $T_1$
- ▶ Snapshot that does not reflect writes by  $T_1$  does not reflect writes by  $T_2$  either

# Consistent snapshots

A read-only transaction observes a **consistent snapshot**:

If  $T_1 \rightarrow T_2$  (e.g.  $T_2$  reads data written by  $T_1$ )...

- ▶ Snapshot reflecting writes by  $T_2$  also reflects writes by  $T_1$
- ▶ Snapshot that does not reflect writes by  $T_1$  does not reflect writes by  $T_2$  either
- ▶ In other words, snapshot is **consistent with causality**
- ▶ Even if read-only transaction runs for a long time

# Consistent snapshots

A read-only transaction observes a **consistent snapshot**:

If  $T_1 \rightarrow T_2$  (e.g.  $T_2$  reads data written by  $T_1$ )...

- ▶ Snapshot reflecting writes by  $T_2$  also reflects writes by  $T_1$
- ▶ Snapshot that does not reflect writes by  $T_1$  does not reflect writes by  $T_2$  either
- ▶ In other words, snapshot is **consistent with causality**
- ▶ Even if read-only transaction runs for a long time

Approach: **multi-version concurrency control** (MVCC)

- ▶ Each read-write transaction  $T_w$  has commit timestamp  $t_w$
- ▶ Every value is tagged with timestamp  $t_w$  of transaction that wrote it (not overwriting previous value)

## Consistent snapshots

A read-only transaction observes a **consistent snapshot**:

If  $T_1 \rightarrow T_2$  (e.g.  $T_2$  reads data written by  $T_1$ )...

- ▶ Snapshot reflecting writes by  $T_2$  also reflects writes by  $T_1$
- ▶ Snapshot that does not reflect writes by  $T_1$  does not reflect writes by  $T_2$  either
- ▶ In other words, snapshot is **consistent with causality**
- ▶ Even if read-only transaction runs for a long time

Approach: **multi-version concurrency control** (MVCC)

- ▶ Each read-write transaction  $T_w$  has commit timestamp  $t_w$
- ▶ Every value is tagged with timestamp  $t_w$  of transaction that wrote it (not overwriting previous value)
- ▶ Read-only transaction  $T_r$  has snapshot timestamp  $t_r$
- ▶  $T_r$  ignores values with  $t_w > t_r$ ; observes most recent value with  $t_w \leq t_r$

# Obtaining commit timestamps

Must ensure that whenever  $T_1 \rightarrow T_2$  we have  $t_1 < t_2$ .

- ▶ Physical clocks may be **inconsistent with causality**

# Obtaining commit timestamps



Must ensure that whenever  $T_1 \rightarrow T_2$  we have  $t_1 < t_2$ .

- ▶ Physical clocks may be **inconsistent with causality**
- ▶ Can we use Lamport clocks instead?
- ▶ Problem: linearizability depends on **real-time order**, and logical clocks may not reflect this!

# Obtaining commit timestamps

Must ensure that whenever  $T_1 \rightarrow T_2$  we have  $t_1 < t_2$ .


- ▶ Physical clocks may be **inconsistent with causality**
- ▶ Can we use Lamport clocks instead?
- ▶ Problem: linearizability depends on **real-time order**, and logical clocks may not reflect this!



# Obtaining commit timestamps

Must ensure that whenever  $T_1 \rightarrow T_2$  we have  $t_1 < t_2$ .

- ▶ Physical clocks may be **inconsistent with causality**
- ▶ Can we use Lamport clocks instead?
- ▶ Problem: linearizability depends on **real-time order**, and logical clocks may not reflect this!







# Obtaining commit timestamps

Must ensure that whenever  $T_1 \rightarrow T_2$  we have  $t_1 < t_2$ .

- ▶ Physical clocks may be **inconsistent with causality**
- ▶ Can we use Lamport clocks instead?
- ▶ Problem: linearizability depends on **real-time order**, and logical clocks may not reflect this!




# TrueTime: explicit physical clock uncertainty



# TrueTime: explicit physical clock uncertainty

Spanner's TrueTime clock returns  $[t_{\text{earliest}}, t_{\text{latest}}]$ .

True physical timestamp must lie within that range.




# TrueTime: explicit physical clock uncertainty

Spanner's TrueTime clock returns  $[t_{\text{earliest}}, t_{\text{latest}}]$ .

True physical timestamp must lie within that range.

On commit, wait for uncertainty  $\delta_i = t_{i,\text{latest}} - t_{i,\text{earliest}}$ .

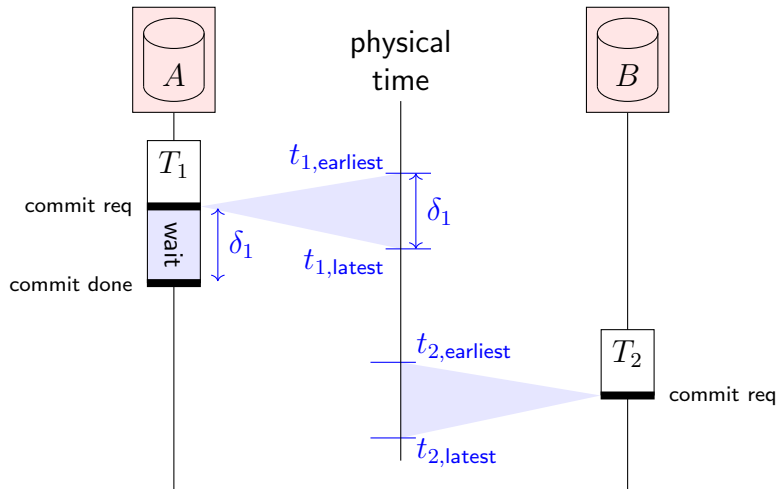


# TrueTime: explicit physical clock uncertainty

Spanner's TrueTime clock returns  $[t_{\text{earliest}}, t_{\text{latest}}]$ .

True physical timestamp must lie within that range.

On commit, wait for uncertainty  $\delta_i = t_{i,\text{latest}} - t_{i,\text{earliest}}$ .

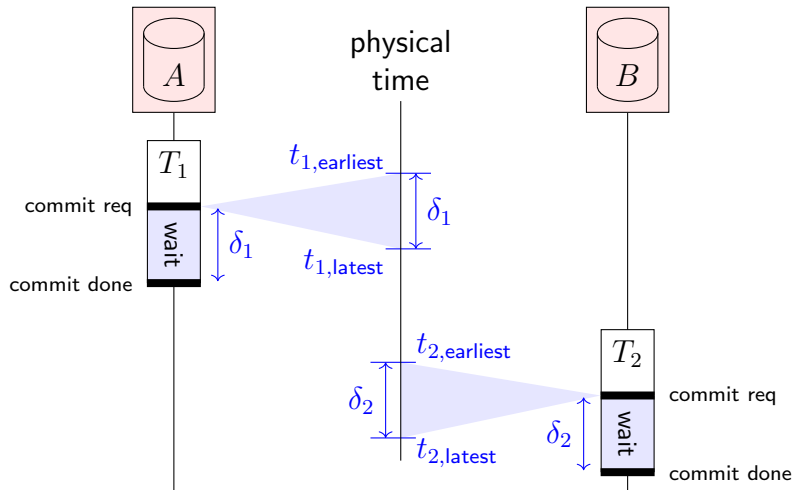


# TrueTime: explicit physical clock uncertainty

Spanner's TrueTime clock returns  $[t_{\text{earliest}}, t_{\text{latest}}]$ .

True physical timestamp must lie within that range.

On commit, wait for uncertainty  $\delta_i = t_{i,\text{latest}} - t_{i,\text{earliest}}$ .




# TrueTime: explicit physical clock uncertainty

Spanner's TrueTime clock returns  $[t_{\text{earliest}}, t_{\text{latest}}]$ .


True physical timestamp must lie within that range.

On commit, wait for uncertainty  $\delta_i = t_{i,\text{latest}} - t_{i,\text{earliest}}$ .



# Determining clock uncertainty in TrueTime

Clock servers with **atomic clock** or **GPS receiver** in each datacenter; servers report their clock uncertainty.






# Determining clock uncertainty in TrueTime


Clock servers with **atomic clock** or **GPS receiver** in each datacenter; servers report their clock uncertainty.

Each node syncs its quartz clock with a server every 30 sec.




# Determining clock uncertainty in TrueTime

Clock servers with **atomic clock** or **GPS receiver** in each datacenter; servers report their clock uncertainty.  
Each node syncs its quartz clock with a server every 30 sec.  
Between syncs, assume worst-case drift of 200ppm.



# Determining clock uncertainty in TrueTime

Clock servers with **atomic clock** or **GPS receiver** in each datacenter; servers report their clock uncertainty.  
Each node syncs its quartz clock with a server every 30 sec.  
Between syncs, assume worst-case drift of 200ppm.



# That's all, folks!

**Any questions?** Email [mk428@cst.cam.ac.uk](mailto:mk428@cst.cam.ac.uk)!

Summary:

- ▶ Distributed systems are everywhere
- ▶ You use them every day: e.g. web apps
- ▶ Key goals: availability, scalability, performance
- ▶ Key problems: concurrency, faults, unbounded latency
- ▶ Key abstractions: replication, broadcast, consensus
- ▶ No one right way, just trade-offs